

# mileage

May 16, 2016

```
In [29]: data = graphlab.SFrame('auto-mpg.csv')
         data
```

Finished parsing file /Users/adityajp/Desktop/ML course/Mileage project/auto-mpg.csv

Parsing completed. Parsed 100 lines in 0.01403 secs.

Finished parsing file /Users/adityajp/Desktop/ML course/Mileage project/auto-mpg.csv

Parsing completed. Parsed 392 lines in 0.014766 secs.

```
-----
Inferred types from first 100 line(s) of file as
column_type_hints=[int,int,float,int,int,float,int,int,str]
If parsing fails due to incorrect types, you can correct
the inferred type list above and pass it to read_csv in
the column_type_hints argument
-----
```

Out[29]: Columns:

```
      mpg      int
cylinders      int
displacement  float
hp            int
weight        int
acceleration  float
model year    int
origin        int
name          str
```

Rows: 392

Data:

mpg	cylinders	displacement	hp	weight	acceleration	model year	origin
18	8	307.0	130	3504	12.0	70	1
15	8	350.0	165	3693	11.5	70	1
18	8	318.0	150	3436	11.0	70	1
16	8	304.0	150	3433	12.0	70	1
17	8	302.0	140	3449	10.5	70	1

15	8	429.0	198	4341	10.0	70	1
14	8	454.0	220	4354	9.0	70	1
14	8	440.0	215	4312	8.5	70	1
14	8	455.0	225	4425	10.0	70	1
15	8	390.0	190	3850	8.5	70	1

name
"chevrolet chevelle malibu"
"buick skylark 320"
"plymouth satellite"
"amc rebel sst"
"ford torino"
"ford galaxie 500"
"chevrolet impala"
"plymouth fury iii"
"pontiac catalina"
"amc ambassador dpl"

[392 rows x 9 columns]

Note: Only the head of the SFrame is printed.

You can use `print_rows(num_rows=m, num_columns=n)` to print more rows and columns.

```
In [2]: import graphlab
```


```
In [128]: def polynomial_sframe(feature, degree, quantity):
    # assume that degree >= 1
    # initialize the SFrame:
    poly_sframe = graphlab.SFrame()
    # and set poly_sframe['power_1'] equal to the passed feature
    poly_sframe[quantity+'power_1'] = feature
    # first check if degree > 1
    if degree > 1:
        # then loop over the remaining degrees:
        # range usually starts at 0 and stops at the endpoint-1. We want it to start at 2 and
        for power in range(2, degree+1):
            # first we'll give the column a name:
            name = quantity + 'power_' + str(power)
            # then assign poly_sframe[name] to the appropriate power of feature
            poly_sframe[name] = feature**power

    return poly_sframe
```

```
In [31]: import matplotlib.pyplot as plt
    %matplotlib inline
```


```
In [32]: plt.plot(data['weight'], data['mpg'], '.')
    #Relation between weight and mpg - Strong correlation
```

```
Out[32]: [<matplotlib.lines.Line2D at 0x117723690>]
```




```
In [33]: plt.plot(data['displacement'],data['mpg'],'.')
          #Relation between displacement and mpg - Time series like correlation
```


```
Out[33]: [<matplotlib.lines.Line2D at 0x118681910>]
```



```
In [34]: plt.plot(data['hp'],data['mpg'],'.')
         #Relation between hp and mpg - Somewhat of a correlation
Out[34]: [<matplotlib.lines.Line2D at 0x118953650>]
```




```
In [38]: plt.plot(data['acceleration'],data['mpg'],'.')
         #Relation between acceleration and mpg - weak correlation
Out[38]: [<matplotlib.lines.Line2D at 0x118f04b90>]
```




```
In [40]: plt.plot(data['model year'],data['mpg'],'.')
         #Relation between model year and mpg - Weak correlation
```

```
Out[40]: [<matplotlib.lines.Line2D at 0x11914ce90>]
```



```
In [43]: plt.plot(data['origin'],data['mpg'],'.')
         #Relation between origin and mpg - No correlation. Discard this information.
```

```
Out[43]: [<matplotlib.lines.Line2D at 0x1199d6a10>]
```




```
In [105]: (training_and_validation, testing) = data.random_split(0.9,seed=0)
          (training, validation) = data.random_split(0.5,seed=0)
```

```
In [180]: #Data to use weight, displacement, hp, acceleration.
```

```
poly_data = polynomial_sframe(data['hp'],15,"hp")
poly_data_weight = polynomial_sframe(data['weight'],15,"weight")
poly_data_disp = polynomial_sframe(data['displacement'],15,"displacement")
poly_data.add_columns(poly_data_weight)
poly_data.add_columns(poly_data_disp)
my_features = poly_data.column_names()
poly_data['mpg'] = data['mpg']
model = graphlab.linear_regression.create(poly_data, target = 'mpg', features = my_features, v
```

```
In [181]: plt.plot(poly_data['hppower_1'],poly_data['mpg'],'.',
                  poly_data['hppower_1'], model.predict(poly_data),'.')
          
```

```
Out[181]: [<matplotlib.lines.Line2D at 0x1206fdf90>,
          <matplotlib.lines.Line2D at 0x120738250>]
```



```
In [177]: from matplotlib import pyplot
import pylab
from mpl_toolkits.mplot3d import Axes3D
import random
import numpy as np

fig = pylab.figure()
fig.set_size_inches(25.5, 14.5)

ax = fig.add_subplot(110, projection='3d')

hp = np.array(poly_data['hpower_1'])
weight = np.array(poly_data['weightpower_1'])
mpg = np.array(poly_data['mpg'])
x_vals = hp
y_vals = weight
z_vals = mpg


pred_x_vals = hp
pred_y_vals = weight
pred_z_vals = model.predict(poly_data)
#random.shuffle(sequence_containing_x_vals)
#random.shuffle(sequence_containing_y_vals)
#random.shuffle(sequence_containing_z_vals)

ax.scatter(x_vals, y_vals, z_vals, c='r', marker='o')
ax.scatter(pred_x_vals, pred_y_vals, pred_z_vals, c='b', marker='^')
ax.set_xlabel('HP')
```

```

ax.set_ylabel('Weight')
ax.set_zlabel('mpg')
pyplot.show()
fig.savefig('test2png.png', dpi=100)

```



```
In [287]: rss_list = []
```

```
#going to consider polynomial upto the power of 15
```

```
for i in range(1,16):
```

```
#Gathering necessary data from the training set
```

```
poly_data = polynomial_sframe(training['hp'],i,"hp")
```

```
poly_data.add_columns(polynomial_sframe(training['weight'],i,"weight"))
```

```
poly_data.add_columns(polynomial_sframe(training['displacement'],i,"displacement"))
```

```
poly_data.add_columns(polynomial_sframe(training['acceleration'],i,"acceleration"))
```

```
#column_names
```

```
my_features = poly_data.column_names()
```

```
poly_data['mpg'] = training['mpg']
```

```
#Creating a regression model
```

```
model = graphlab.linear_regression.create(poly_data, target = 'mpg', features = my_features)
```

```
#Getting necessary data from the validation set
```

```
validation_data = polynomial_sframe(validation['hp'],i,"hp")
```

```
validation_data.add_columns(polynomial_sframe(validation['weight'],i,"weight"))
```

```
validation_data.add_columns(polynomial_sframe(validation['displacement'],i,"displacement"))
```

```
validation_data.add_columns(polynomial_sframe(validation['acceleration'],i,"acceleration"))
```

```
#Predicted mpg
```



```

prediction = model.predict(validation_data)

#Computing the RSS (Residual sum of squares)
difference = validation['mpg'] - prediction
difference = difference ** 2 #Squaring all elements in the SFrame

difference = np.array(difference)
rss = np.sum(difference)

rss_list.append((rss,i)) #Append degree of the polynomial and RSS (Used for sorting later)

sorted_rss = sorted(rss_list, key=lambda tup: tup[0])
degree = sorted_rss[0][1]

for i in sorted_rss:
    print i
    #RSS is least for polynomial of degree 7 and highest for polynomial of degree 1 and RSS of de

(3221.5765214009725, 3)
(3226.0471238434775, 2)
(3306.1360244889956, 4)
(3345.6813203488928, 5)
(3403.116566599871, 6)
(3627.8744744869477, 7)
(3822.5616595836882, 1)
(4025.5880719590386, 8)
(4540.5537163976205, 9)
(5148.9033698786106, 10)
(5842.5112581298999, 11)
(6601.1162841234327, 12)
(7389.8261699294908, 13)
(8172.6708302590287, 14)
(8922.7567327277957, 15)

In [288]: #Gathering testing data into an SFrame
def prediction(testing, model):
    testing_data = polynomial_sframe(testing['hp'],degree,"hp")
    testing_data.add_columns(polynomial_sframe(testing['weight'],degree,"weight"))
    testing_data.add_columns(polynomial_sframe(testing['displacement'],degree,"displacement"))
    testing_data.add_columns(polynomial_sframe(testing['acceleration'],degree,"acceleration"))
    #Prediction step
    testing_prediction = model.predict(testing_data)

    return (testing_prediction)

#Computing the RSS (Residual sum of squares)
predicted['mpg'] = testing['mpg']
predicted['predicted_results'] = prediction(testing,model)

difference = testing['mpg'] - predicted['predicted_results']
predicted['difference'] = difference
difference = difference ** 2 #Squaring all elements in the SFrame

difference = np.array(difference)

```

```
rss = np.sum(difference)
```

```
print rss
```

```
5631.4898989
```

```
In [285]: #predicted['mpg','predicted','difference'].print_rows(100,100)
```

```
print predicted
```

mpg	predicted	predicted_results	differnce	difference
15	-24.5863221627	41.936200147	-26.936200147	-26.936200147
22	27.7579668808	17.6166743233	4.38332567671	4.38332567671
28	35.9924972688	21.9232293157	6.07677068425	6.07677068425
24	35.6695108685	21.5328612405	2.46713875948	2.46713875948
13	-6.2817479709	26.3984305342	-13.3984305342	-13.3984305342
13	0.156837826423	22.8234156944	-9.82341569445	-9.82341569445
14	0.121701026037	21.5109049503	-7.51090495033	-7.51090495033
22	35.6418542922	21.117153017	0.882846982955	0.882846982955
20	26.4984432094	17.5812026779	2.41879732208	2.41879732208
29	38.8228638577	23.155953446	5.84404655405	5.84404655405

```
[37 rows x 5 columns]
```

Note: Only the head of the SFrame is printed.

You can use `print_rows(num_rows=m, num_columns=n)` to print more rows and columns.

```
In [ ]:
```